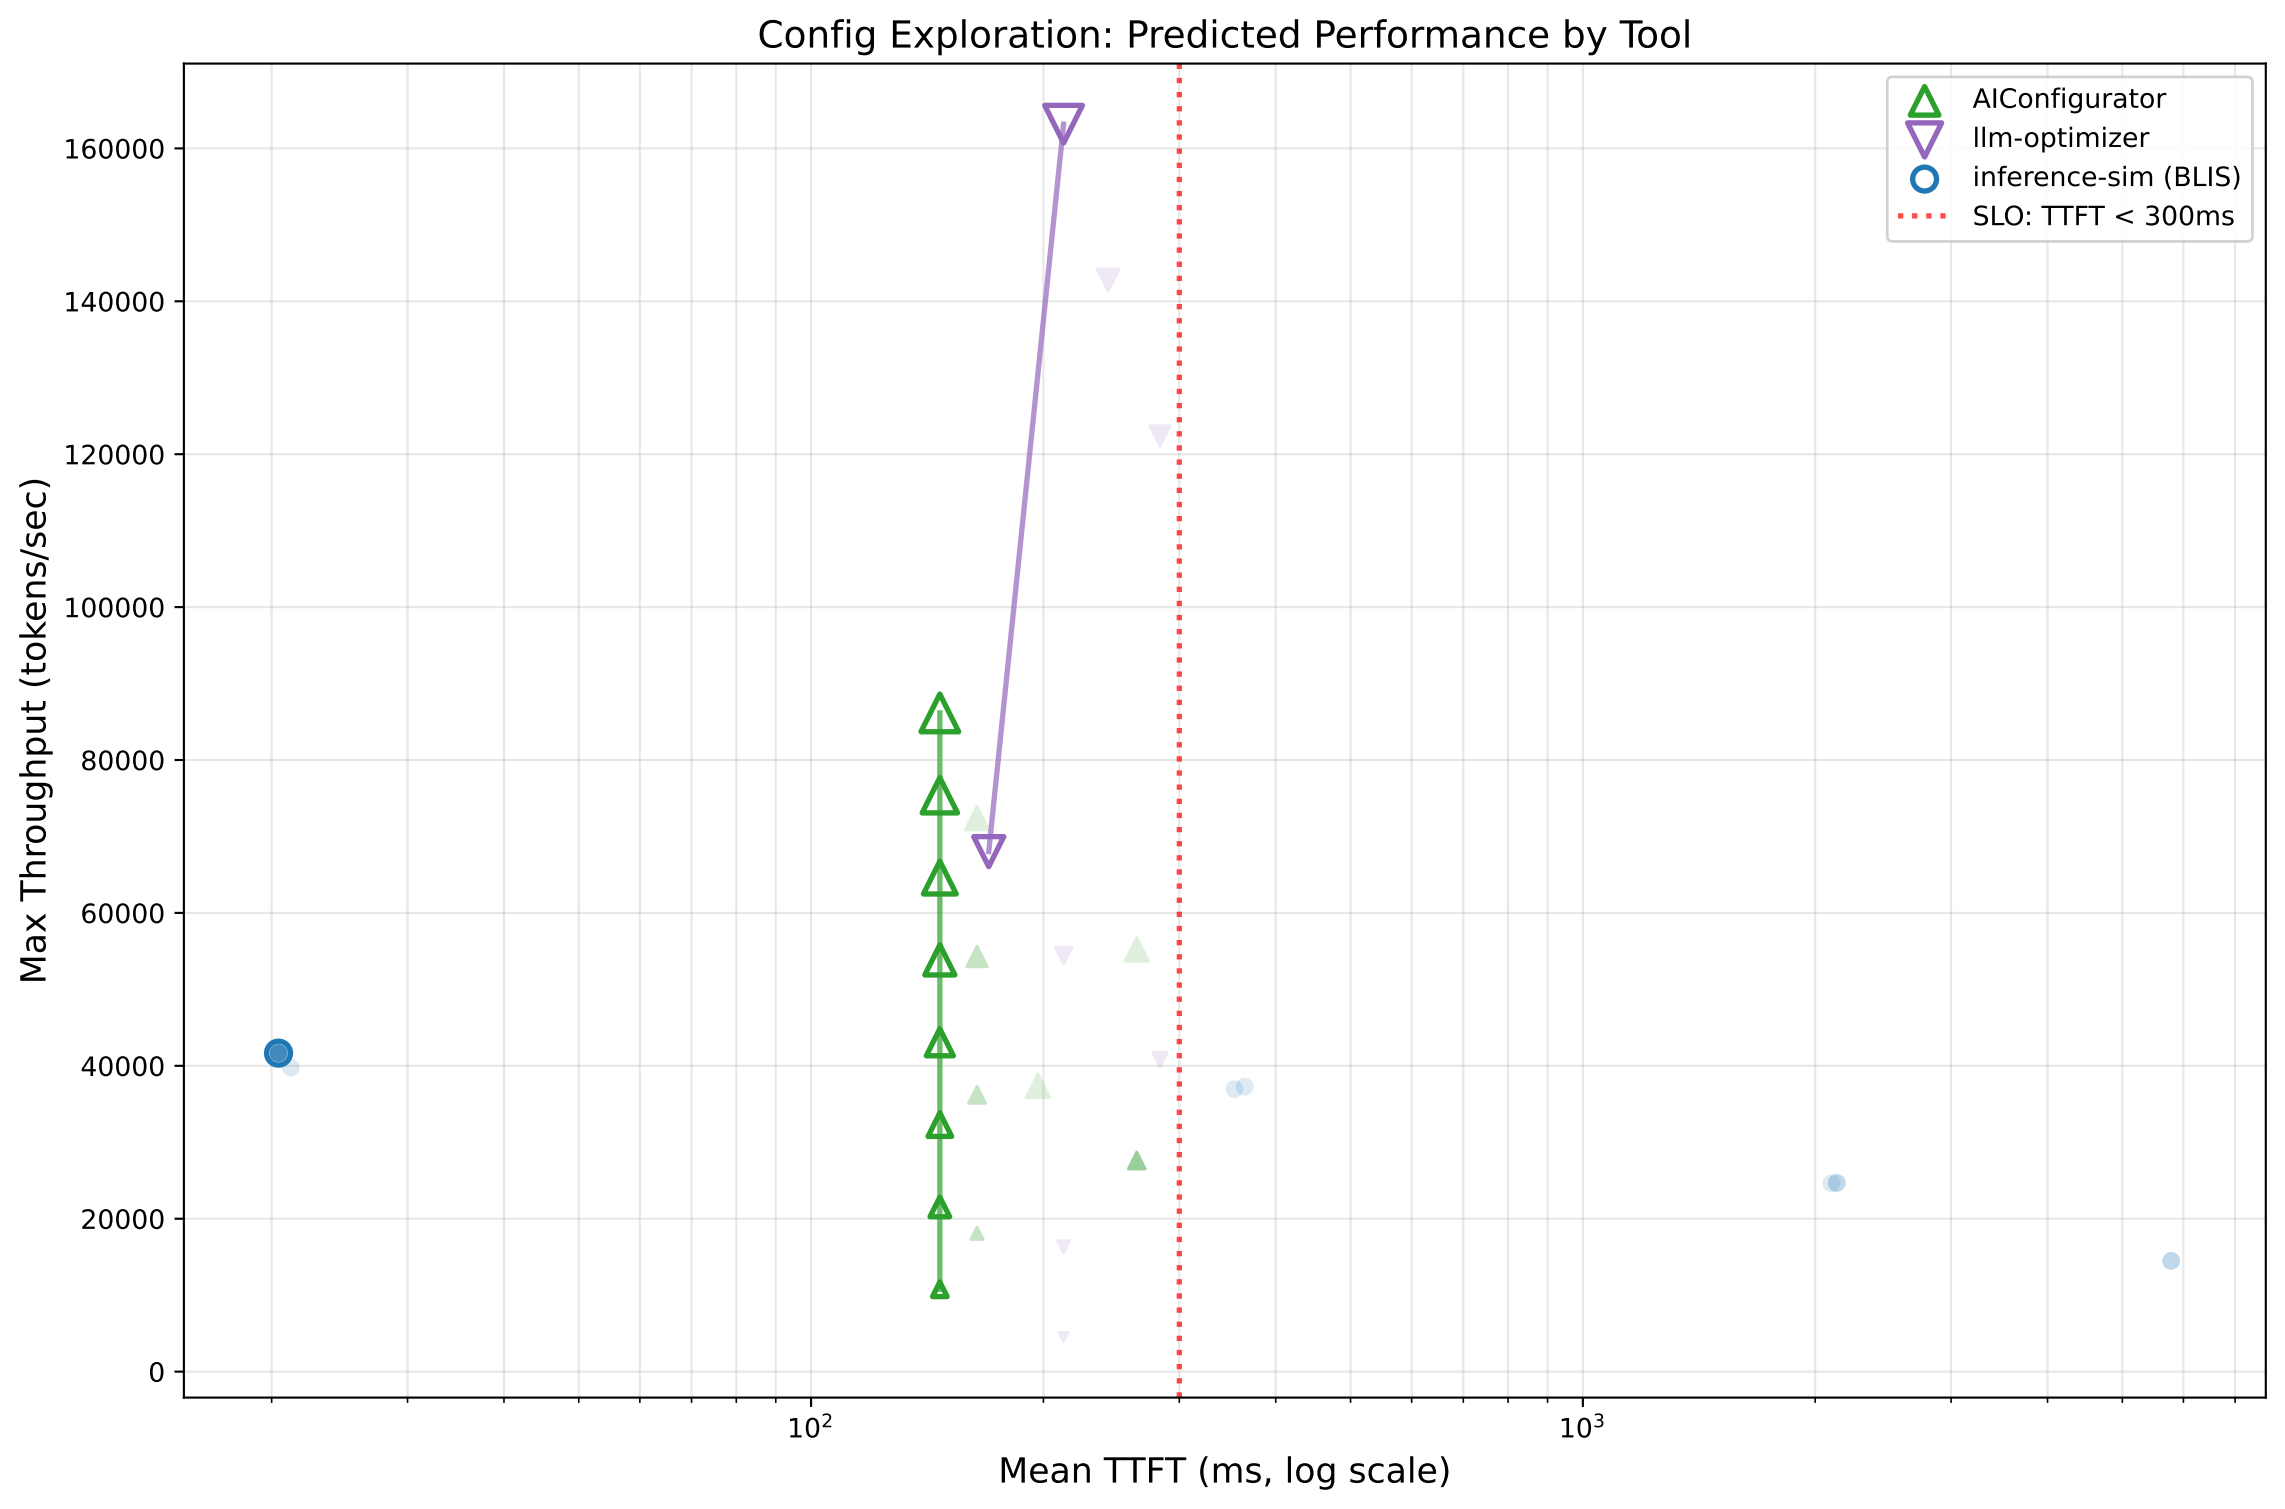
Config Exploration: Predicted Performance by Tool
AIConfigurator
llm-optimizer
inference-sim (BLIS)
SLO: TTFT < 300ms
Max Throughput (tokens/sec)
160000
140000
120000
100000
80000
60000
40000
20000
0
10^2
10^3
Mean TTFT (ms, log scale)