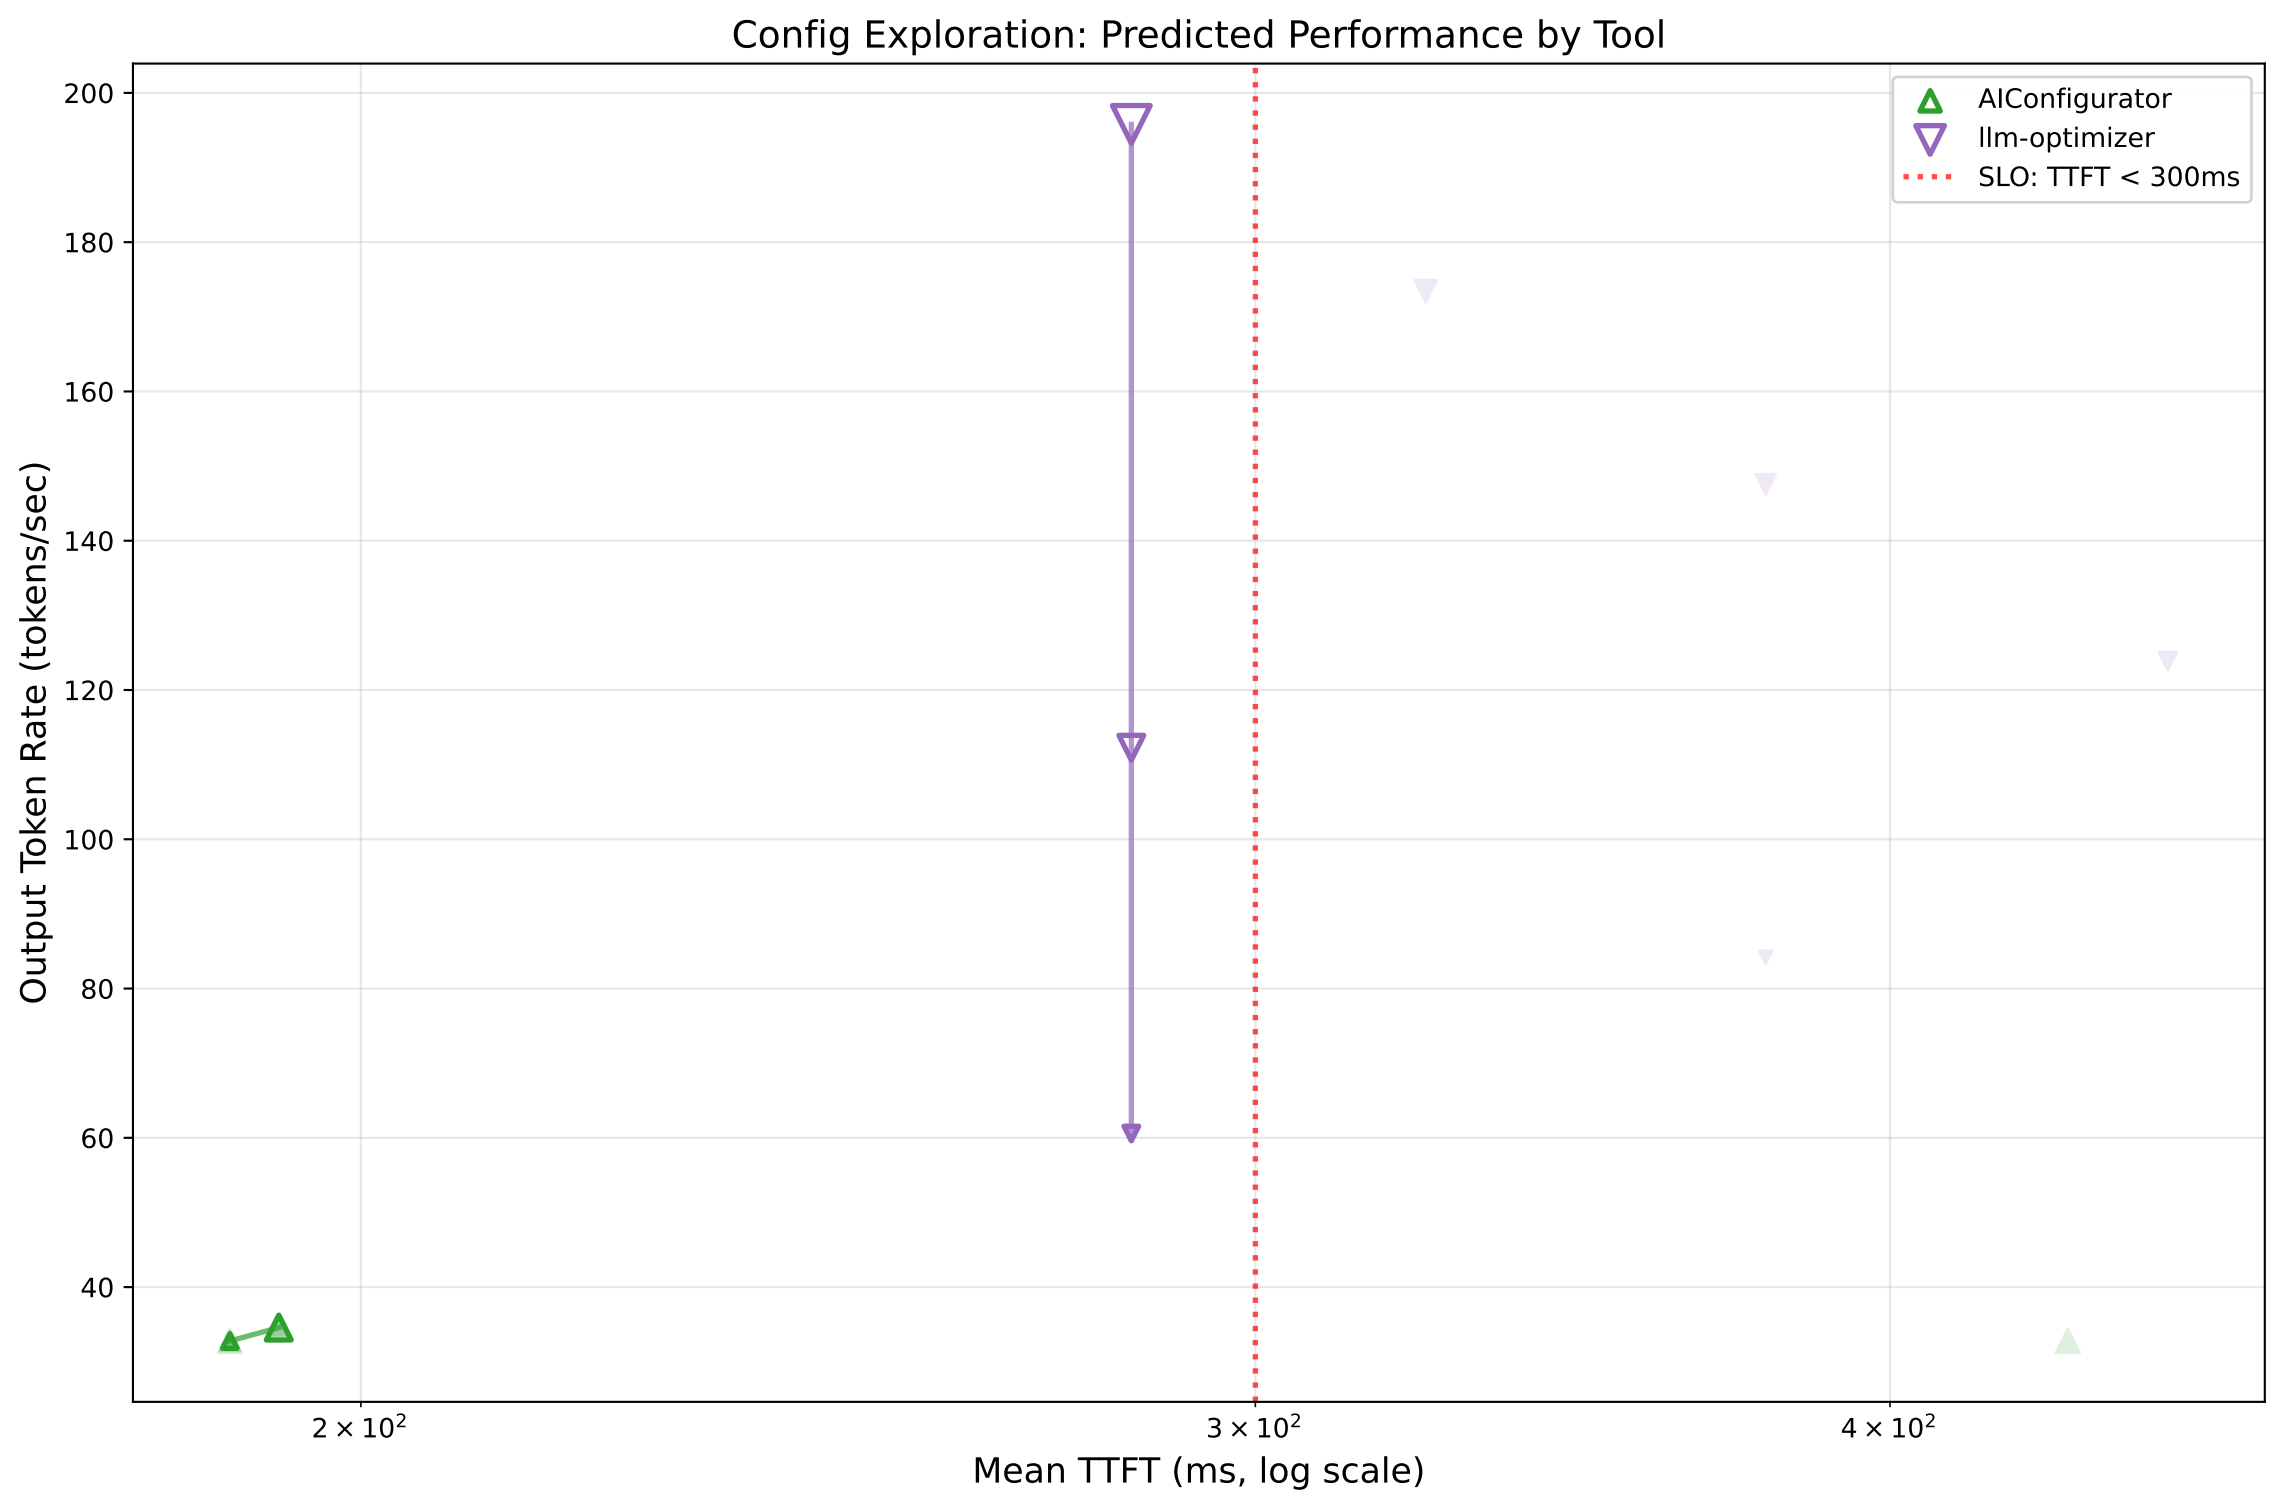
Config Exploration: Predicted Performance by Tool
AIConfigurator
llm-optimizer
SLO: TTFT < 300ms
Output Token Rate (tokens/sec)
200
180
160
140
120
100
80
60
40
2×10^2
3×10^2
4×10^2
Mean TTFT (ms, log scale)