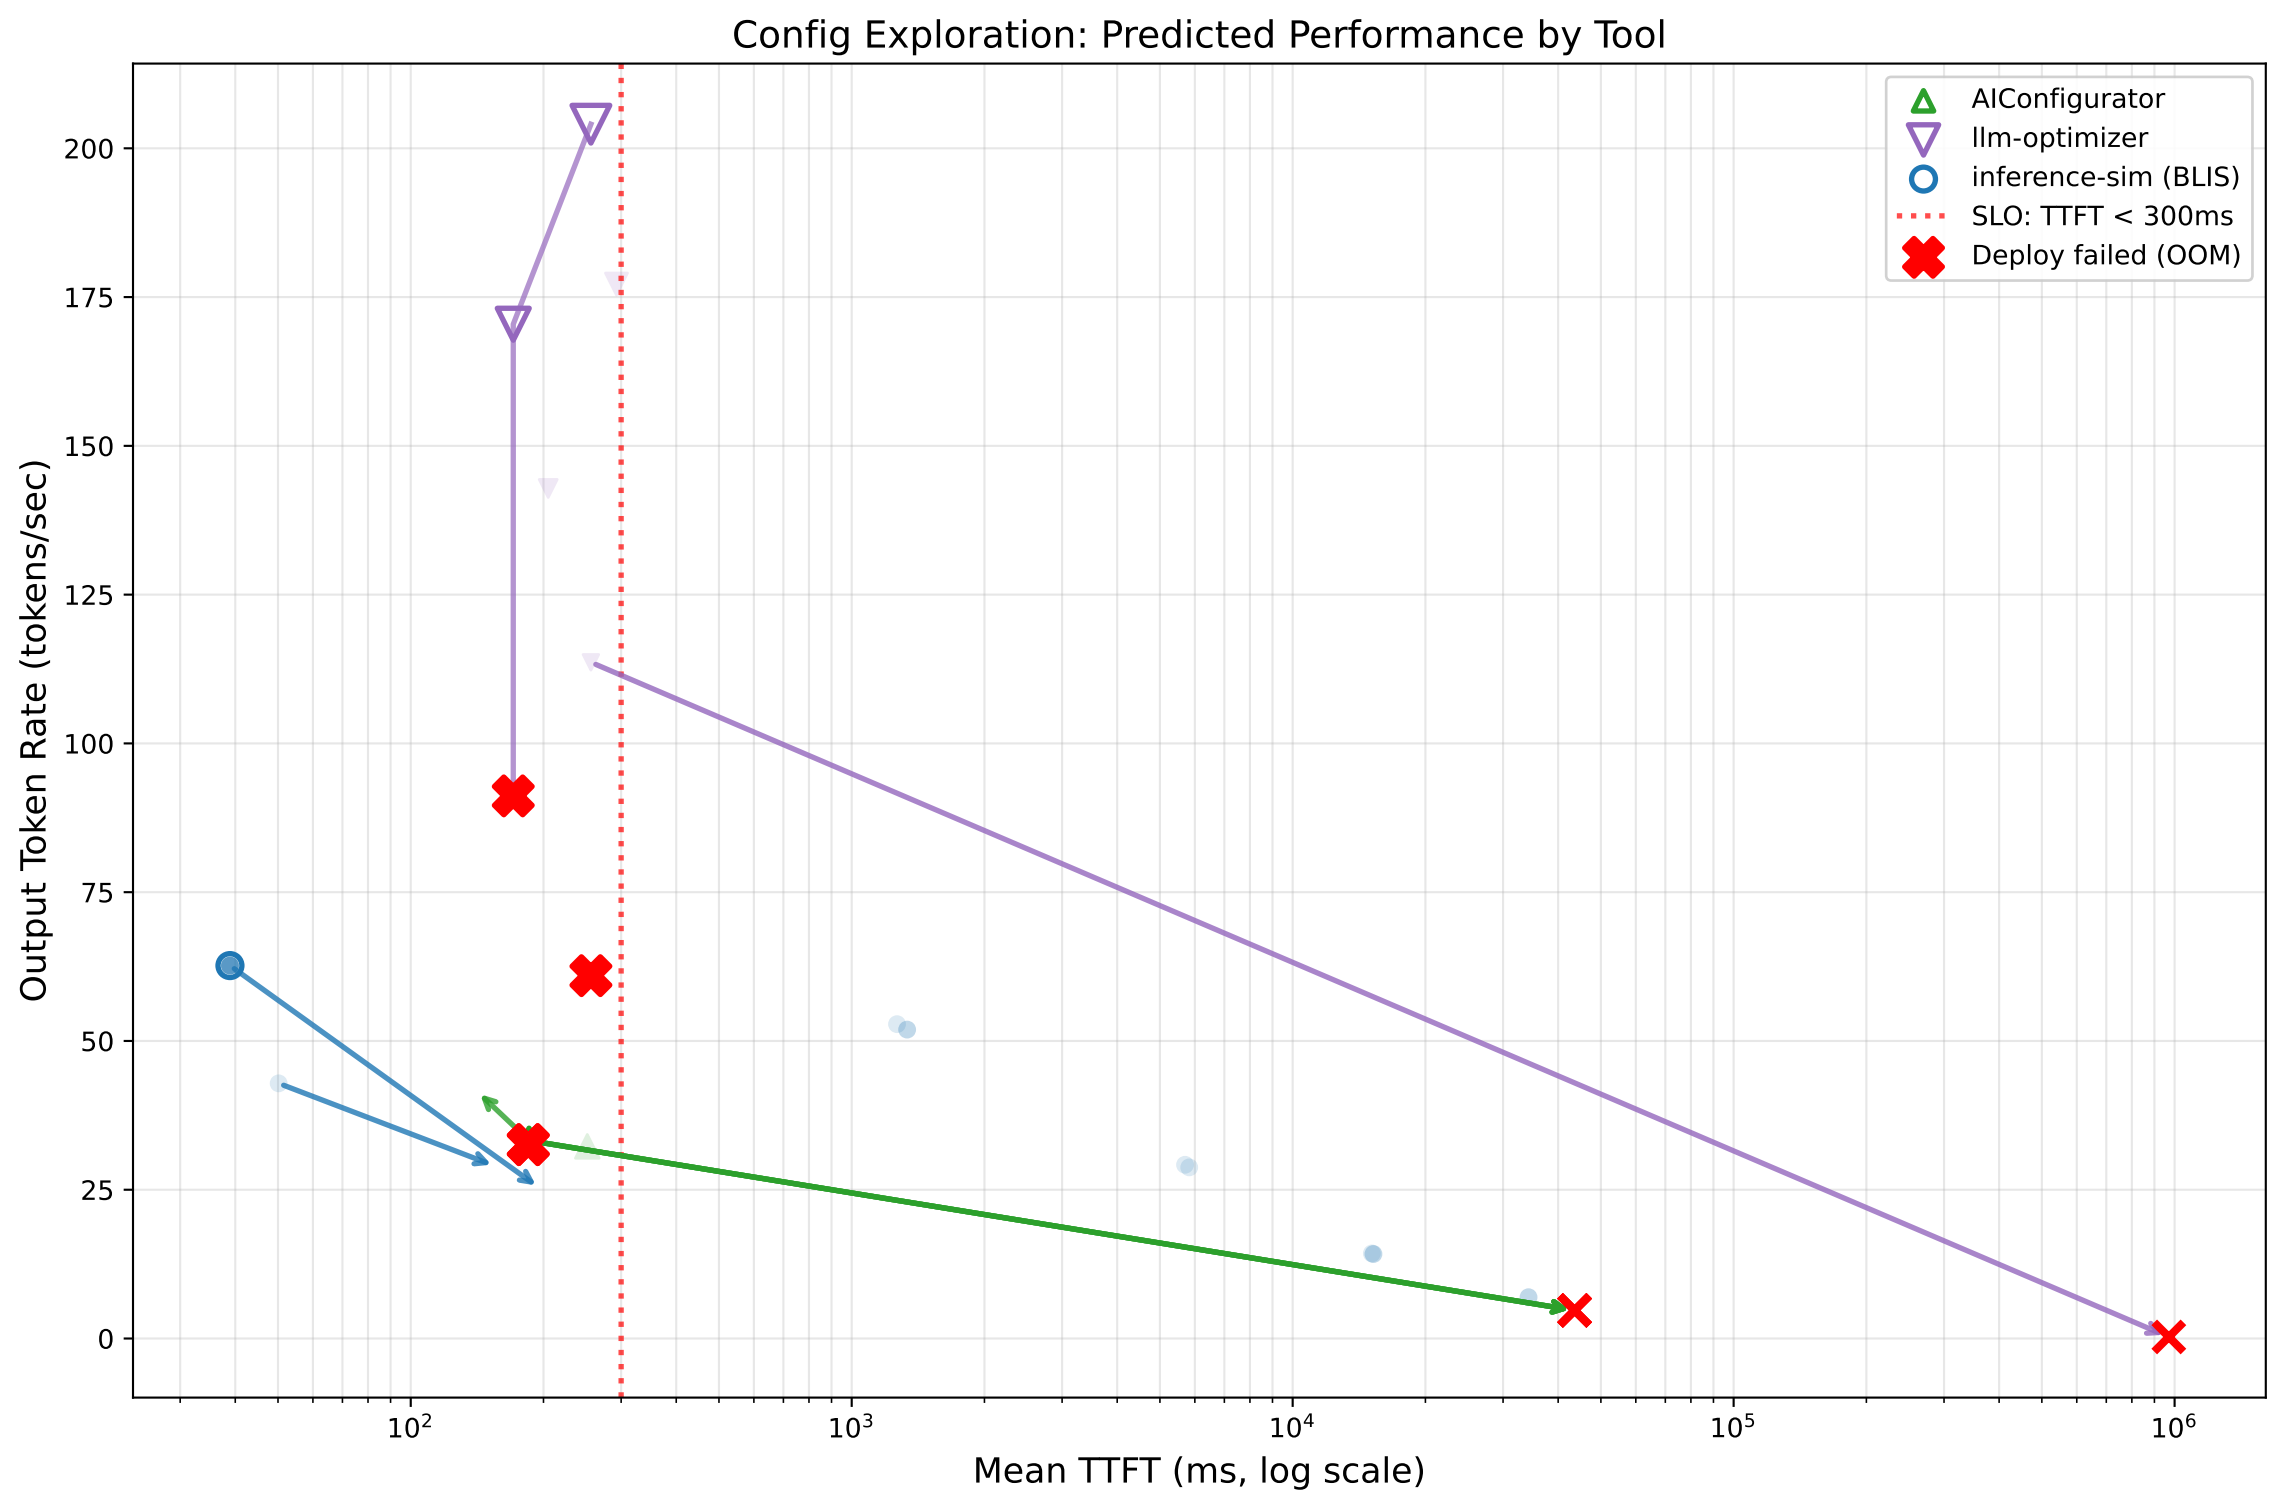
Config Exploration: Predicted Performance by Tool
AIConfigurator
llm-optimizer
inference-sim (BLIS)
SLO: TTFT < 300ms
Deploy failed (OOM)
Output Token Rate (tokens/sec)
Mean TTFT (ms, log scale)
0
25
50
75
100
125
150
175
200
10^2
10^3
10^4
10^5
10^6