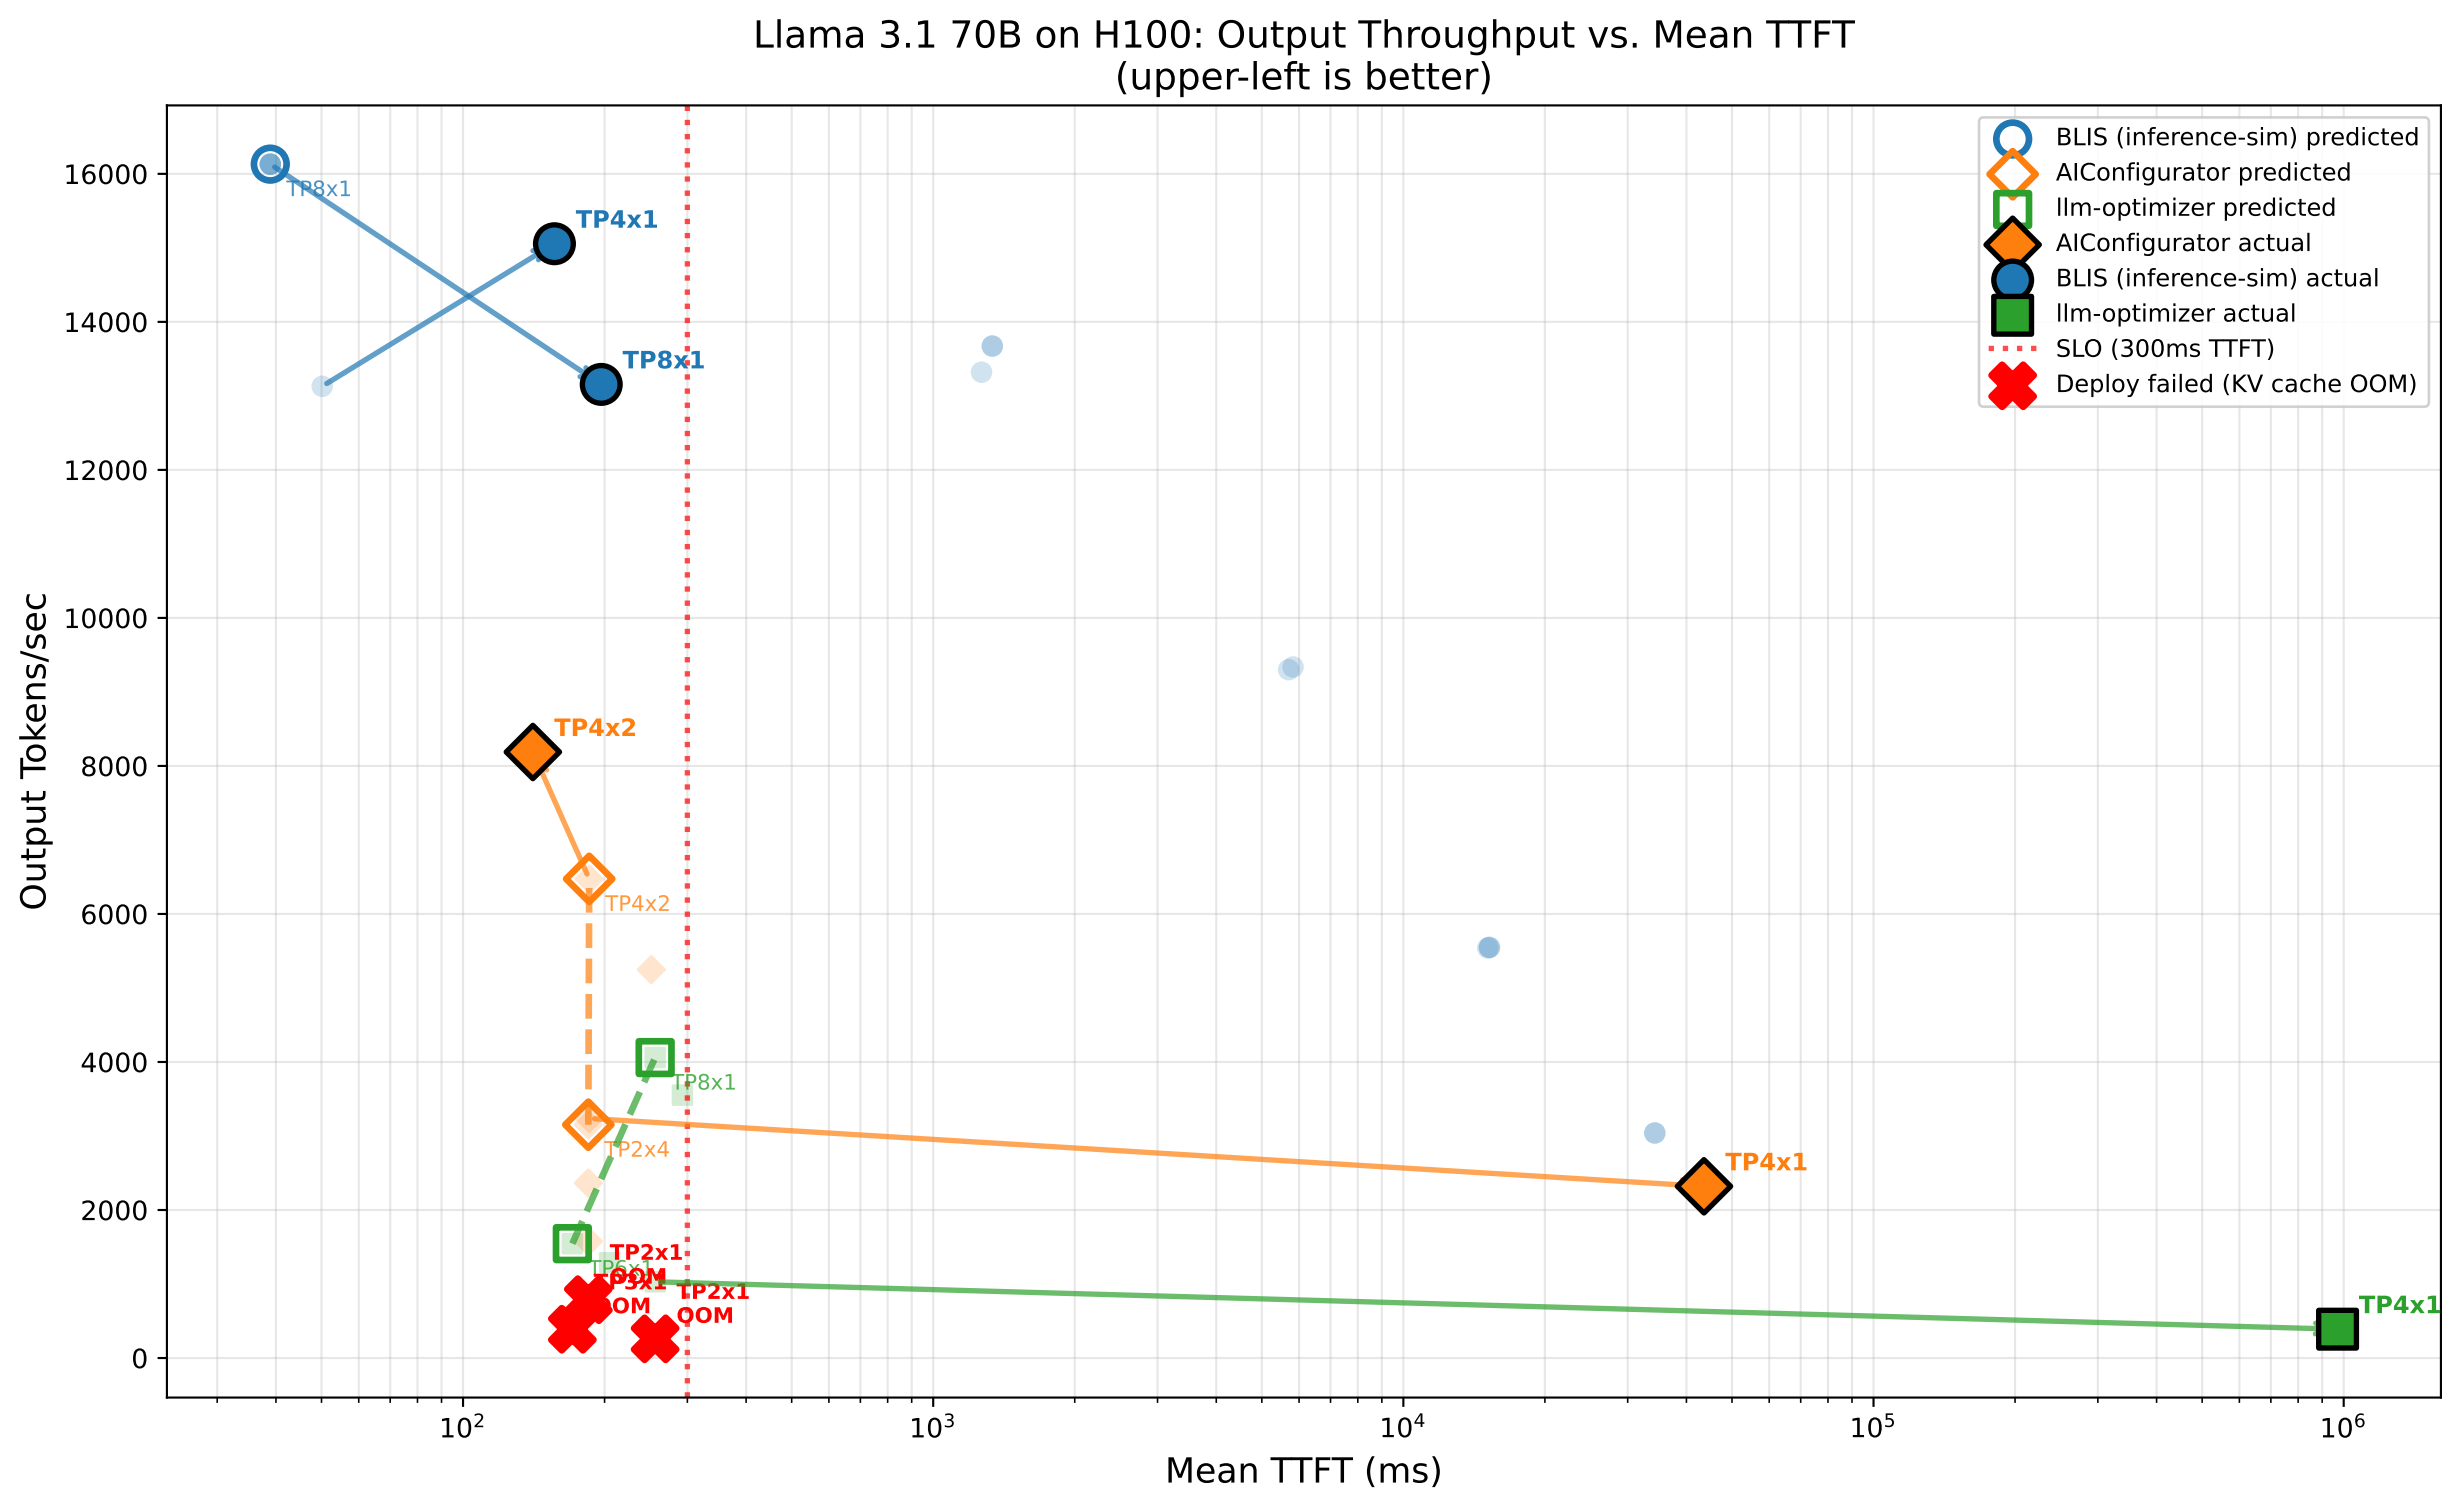
Llama 3.1 70B on H100: Output Throughput vs. Mean TTFT
(upper-left is better)
Output Tokens/sec
Mean TTFT (ms)
BLIS (inference-sim) predicted
AIConfigurator predicted
llm-optimizer predicted
AIConfigurator actual
BLIS (inference-sim) actual
llm-optimizer actual
SLO (300ms TTFT)
Deploy failed (KV cache OOM)
TP8x1
TP4x1
TP8x1
TP4x2
TP4x2
TP8x1
TP2x4
TP4x1
TP2x1
OOM
TP2x1
OOM
TP4x1