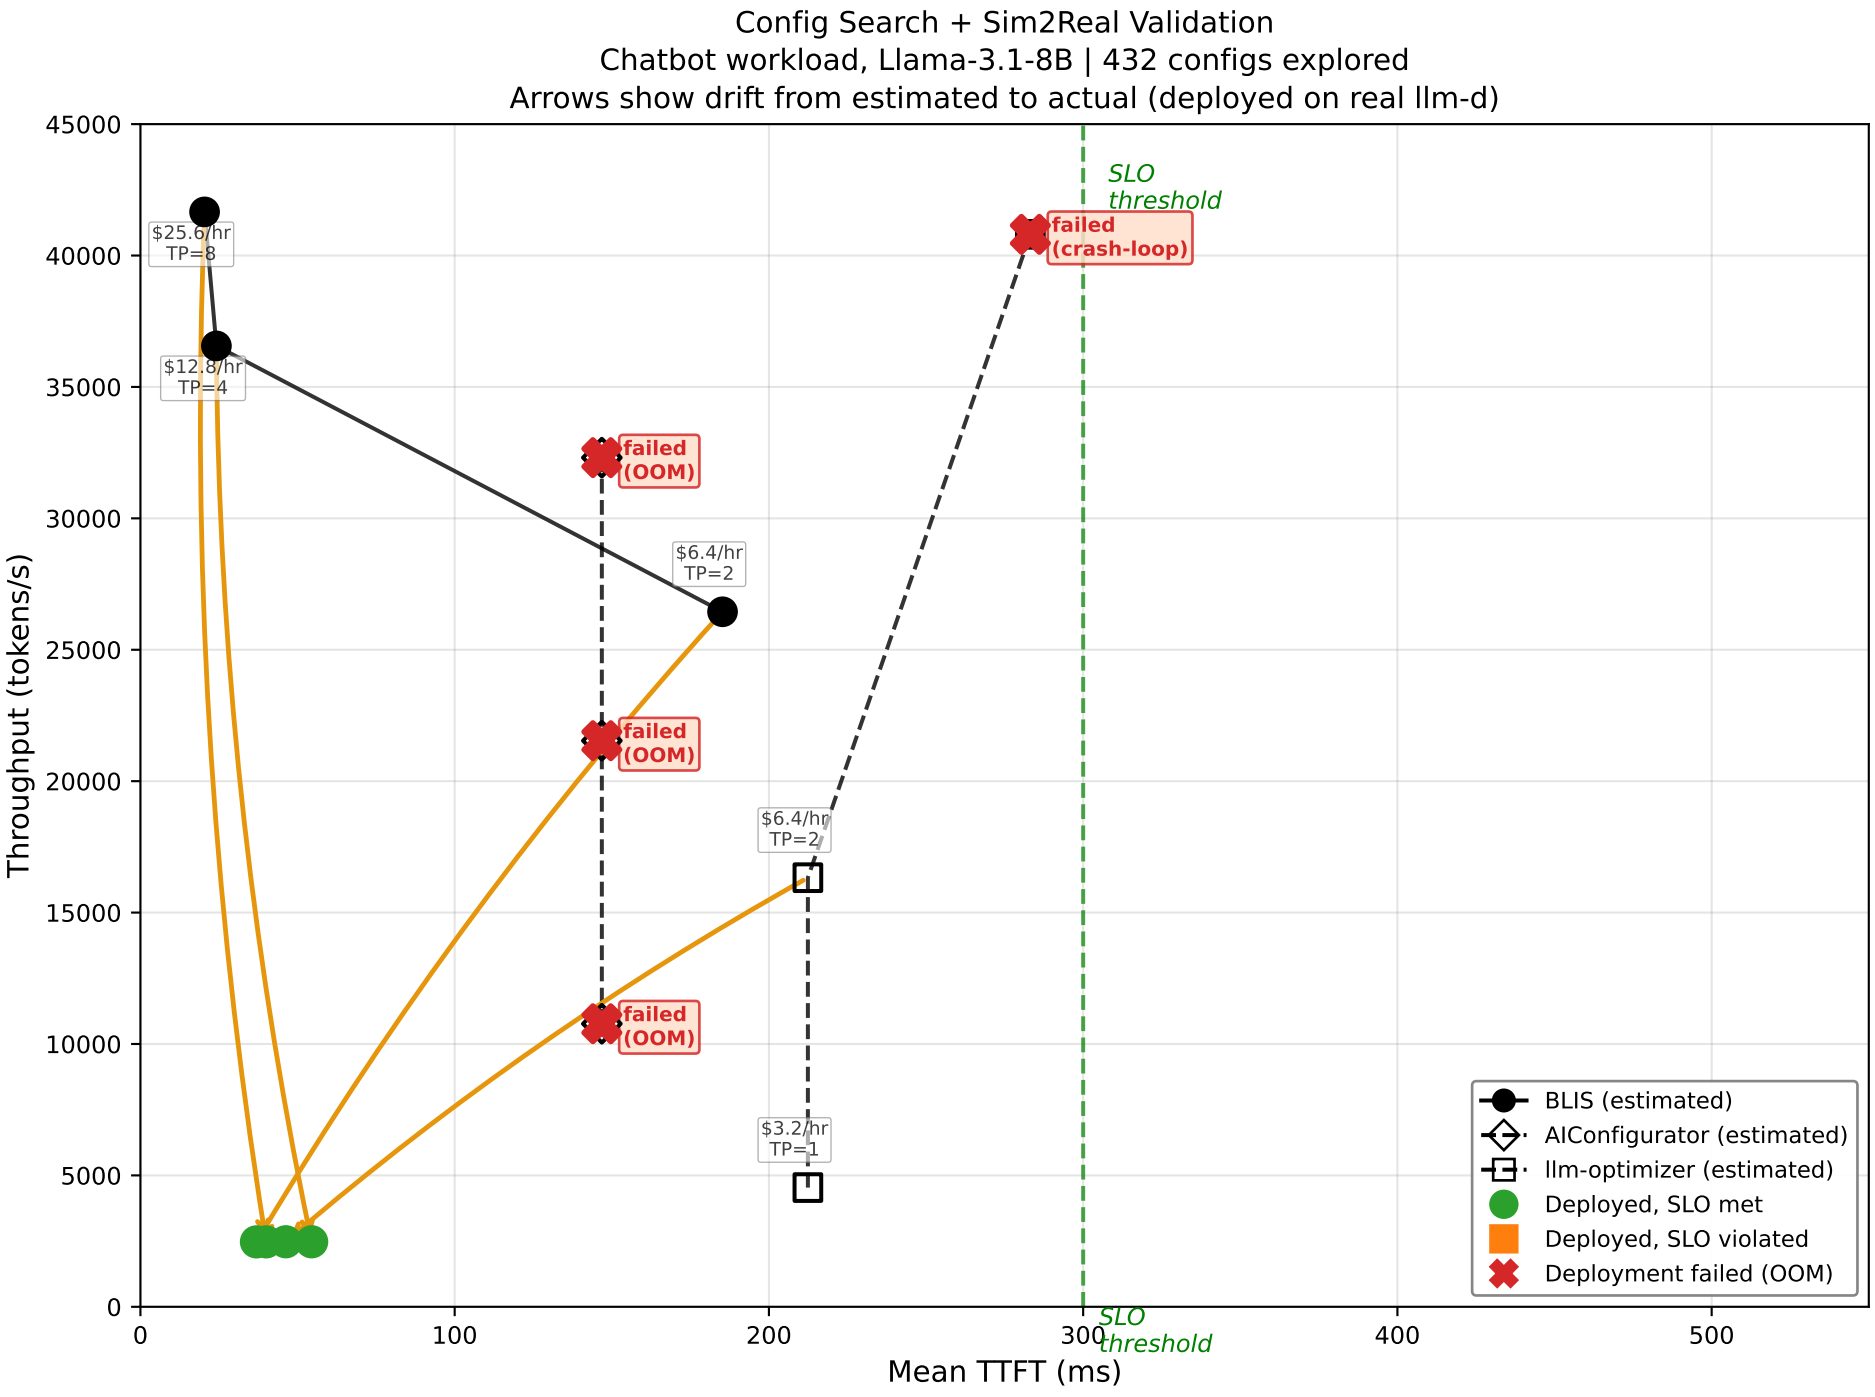
Config Search + Sim2Real Validation
Chatbot workload, Llama-3.1-8B | 432 configs explored
Arrows show drift from estimated to actual (deployed on real llm-d)
Throughput (tokens/s)
Mean TTFT (ms)
SLO threshold
SLO threshold
$25.6/hr TP=8
$12.8/hr TP=4
$6.4/hr TP=2
$6.4/hr TP=2
$3.2/hr TP=1
failed (crash-loop)
failed (OOM)
failed (OOM)
failed (OOM)
BLIS (estimated)
AIConfigurator (estimated)
llm-optimizer (estimated)
Deployed, SLO met
Deployed, SLO violated
Deployment failed (OOM)