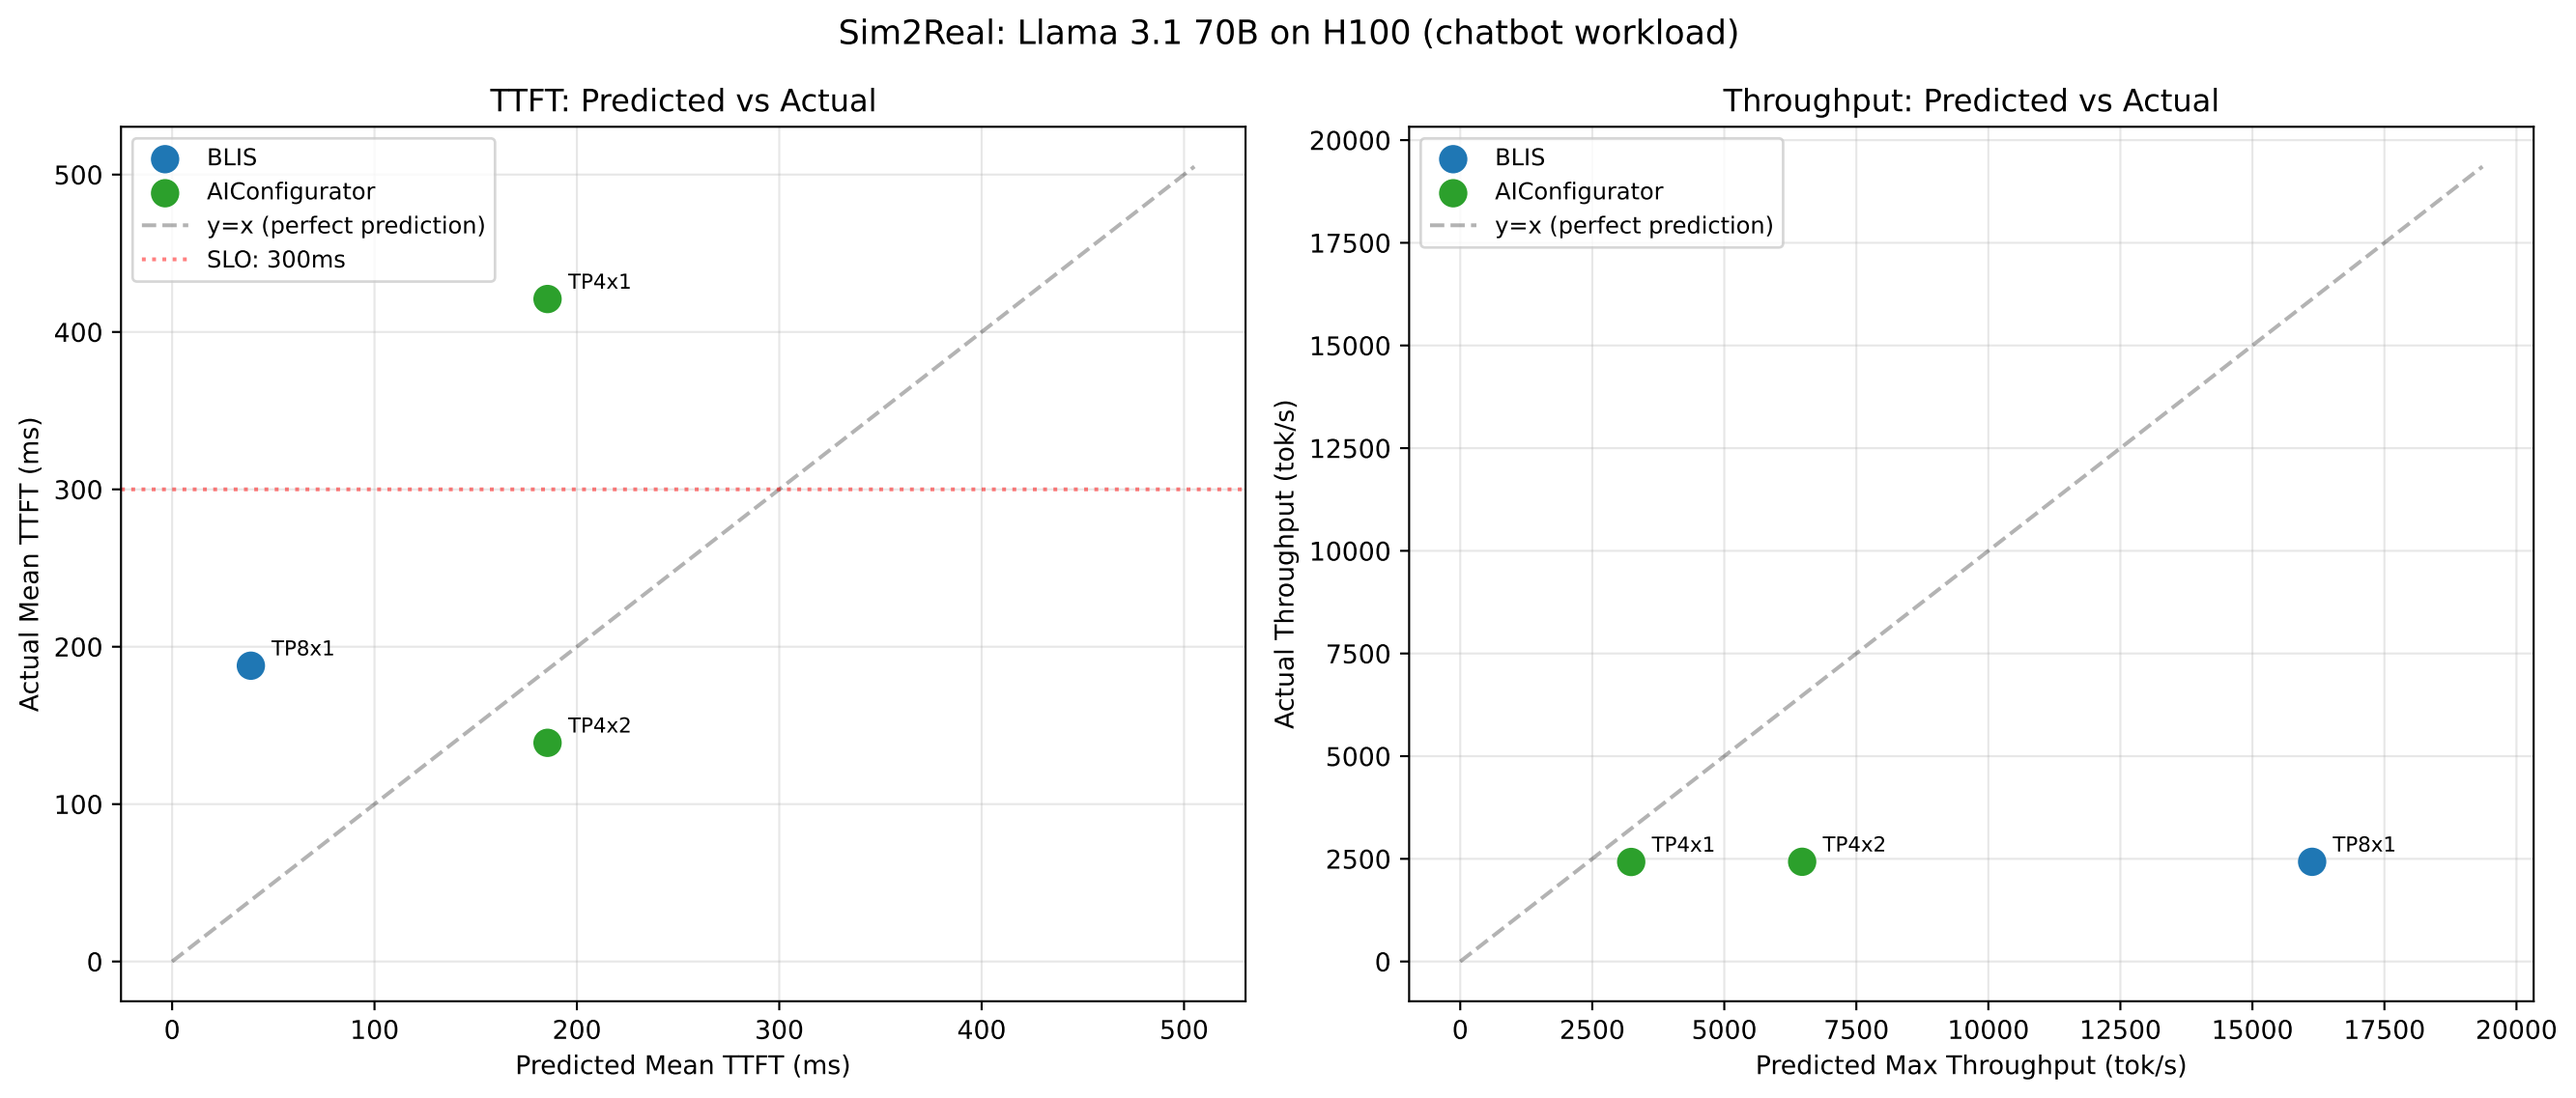
Sim2Real: Llama 3.1 70B on H100 (chatbot workload)
TTFT: Predicted vs Actual
BLIS
AIConfigurator
y=x (perfect prediction)
SLO: 300ms
TP4x1
TP8x1
TP4x2
Actual Mean TTFT (ms)
Predicted Mean TTFT (ms)
0
100
200
300
400
500
Throughput: Predicted vs Actual
BLIS
AIConfigurator
y=x (perfect prediction)
TP4x1
TP4x2
TP8x1
Actual Throughput (tok/s)
Predicted Max Throughput (tok/s)
0
2500
5000
7500
10000
12500
15000
17500
20000