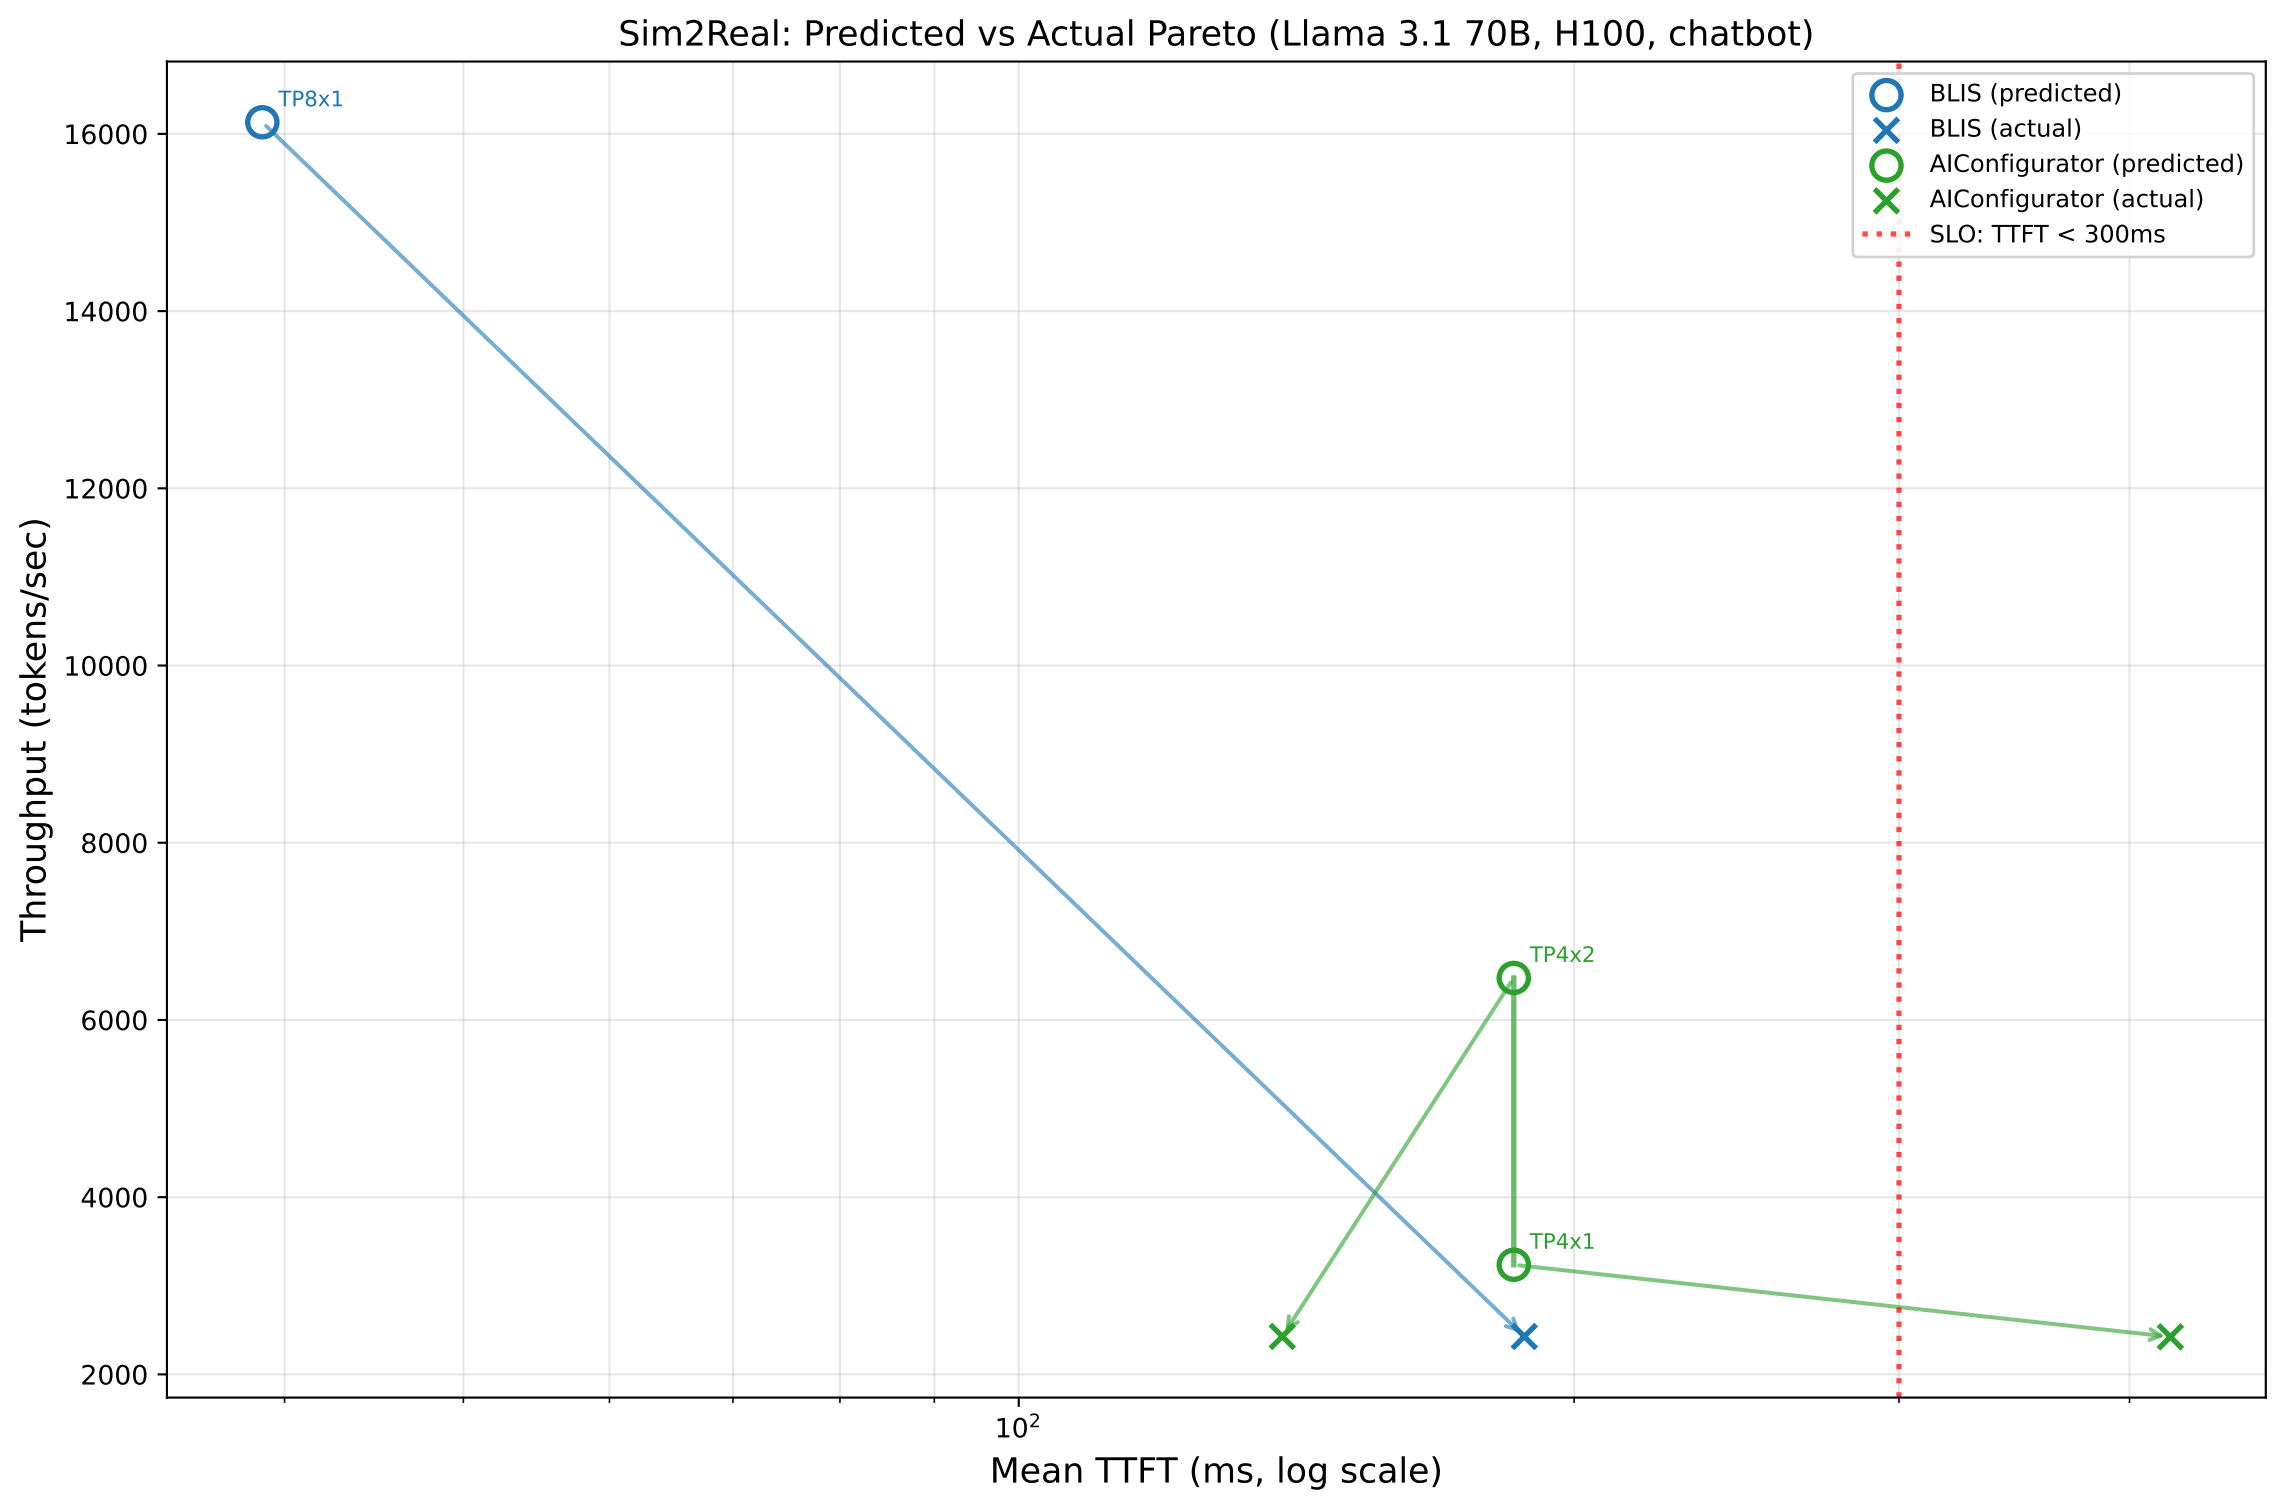
Sim2Real: Predicted vs Actual Pareto (Llama 3.1 70B, H100, chatbot)
BLIS (predicted)
BLIS (actual)
AIConfigurator (predicted)
AIConfigurator (actual)
SLO: TTFT < 300ms
TP8x1
TP4x2
TP4x1
Throughput (tokens/sec)
Mean TTFT (ms, log scale)
16000
14000
12000
10000
8000
6000
4000
2000
$10^2$